Additional References

Introductory level:

"The Book of Why", Judea Pearl and Dana Mackenzie.

Describes the historical developments – from probabilistic machine learning to causal data science – in a very engaging manner. Last chapter discusses the importance of causality for artificial intelligence in general.

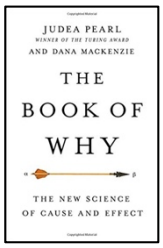

"Causal Inference in Statistics: A Primer", Judea Pearl, Madelyn Glymour, and Nicholas P. Jewell

Probably most accessible introduction into DAG methodology. Technical compendium to the Book of Why.

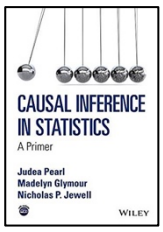

"Counterfactuals and Causal Inference", Stephen L. Morgan and Christopher Whinship

Intuitive introduction into some of the most prominent causal inference techniques, with a discussion of backdoor adjustment and DAGs

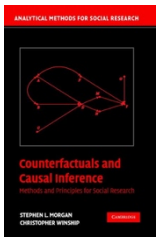

"Causal Inference: The Mixtape", Scott Cunningham.

Discussion of backdoor adjustment and collider bias from an econ perspective. Free download: https://www.scunning.com/mixtape.html

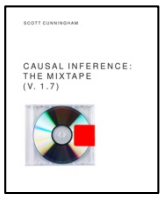

More advanced level:

"Causality: Models, Reasoning, and Inference", Judea Pearl. 2nd edition, 2009.

The causal inference "bible". Very comprehensive, but can be hard to read at times.

"Elements of Causal Inference: Foundations and Learning Algorithms", Jonas Peters, Dominik Janzing and Bernhard Schölkopf.

Probably most comprehensive textbook on causal discovery techniques. Free download: https://mitpress.mit.edu/books/elements-causal-inference

"Actual Causality", Joseph Y. Halpern

Axiomatic foundations of structural causal models and counterfactual reasoning.

"Making Things Happen", James Woodward

Philosophical considerations around a manipulability theory of causation.